

Balancing Search Trees

- Balancing Binary Search Trees
- Operations for Rebalancing
- Tree Rotation
- Insertion at Root
- Tree Partitioning
- Periodic Rebalancing
- Randomised BST Insertion
- An Application of BSTs: Sets

❖ Balancing Binary Search Trees

Observation: order of insertion into a tree affects its height

- **worst case**: keys inserted in ascending/descending order
(effectively have a linked list, so search cost is $O(n)$)
- **best case** (for at-leaf insertion): keys inserted in pre-order
(tree height \Rightarrow search cost is $O(\log n)$; tree is balanced)
- **average case**: keys inserted in **random** order
(tree height \Rightarrow search cost is $O(\log n)$; but cost \geq best case)

Goal: build binary search trees which have

- minimum height \Rightarrow minimum worst case search cost

❖ ... Balancing Binary Search Trees

Perfectly-balanced tree with N nodes has

- \forall nodes, $\text{abs}(\text{\#nodes}(\text{LeftSubtree}) - \text{\#nodes}(\text{RightSubtree})) < 2$
- height of $\log_2 N \Rightarrow$ worst case search $O(\log N)$

Three *strategies* to improving worst case search in BSTs:

- **randomise** — reduce chance of worst-case scenario occurring
- **amortise** — do more work at insertion to make search faster
- **optimise** — implement all operations with performance bounds

❖ Operations for Rebalancing

To assist with rebalancing, we consider new operations:

Left rotation

- move right child to root; rearrange links to retain order

Right rotation

- move left child to root; rearrange links to retain order

Insertion at root


- each new item is added as the new root node

Partition

- rearrange tree around specified node (push it to root)

❖ Tree Rotation

Rotation operations:



Note: tree is ordered, $t_1 < n_2 < t_2 < n_1 < t_3$

❖ ... Tree Rotation

Method for rotating tree T right:

- n_1 is current root; n_2 is root of n_1 's left subtree
- n_1 gets new left subtree, which is n_2 's right subtree
- n_1 becomes root of n_2 's new right subtree
- n_2 becomes new root
- n_2 's left subtree is unchanged


Left rotation: swap left/right in the above.

Rotation requires simple, localised pointer rearrangements

Cost of tree rotation: $O(1)$

❖ ... Tree Rotation

Example of right rotation:



COMP2521 20T1 ♦ Balancing Search Trees ♦ [6/30]

❖ ... Tree Rotation

Algorithm for right rotation:

```
rotateRight( $n_1$ ):  
|   Input   tree  $n_1$   
|   Output  $n_1$  rotated to the right  
|  
|   if  $n_1$  is empty  $\vee$  left( $n_1$ ) is empty then  
|       return  $n_1$   
|   end if  
|    $n_2$  = left( $n_1$ )  
|   left( $n_1$ ) = right( $n_2$ )  
|   right( $n_2$ ) =  $n_1$   
|   return  $n_2$ 
```


❖ ... Tree Rotation

Algorithm for left rotation:

```
rotateLeft( $n_2$ ):  
|   Input   tree  $n_2$   
|   Output  $n_2$  rotated to the left  
|  
|   if  $n_2$  is empty  $\vee$  right( $n_2$ ) is empty then  
|       return  $n_2$   
|   end if  
|    $n_1$ =right( $n_2$ )  
|   right( $n_2$ )=left( $n_1$ )  
|   left( $n_1$ )= $n_2$   
|   return  $n_1$ 
```

❖ ... Tree Rotation

Cost considerations for tree rotation

- the rotation operation is cheap $O(1)$
- if applied appropriately, will tend to improve tree balance

Sometimes rotation is applied from leaf to root, along one branch

- cost of this is $O(\text{height})$
- payoff is improved balance which reduces height
- reduced height pushes search cost towards $O(\log n)$

❖ Insertion at Root

Previous discussion of BSTs did insertion at leaves.

Different approach: insert new item at root.

Potential disadvantages:

- large-scale rearrangement of tree for each insert (apparently)

Potential advantages:

- recently-inserted items are close to root
- lower cost if recent items more likely to be searched

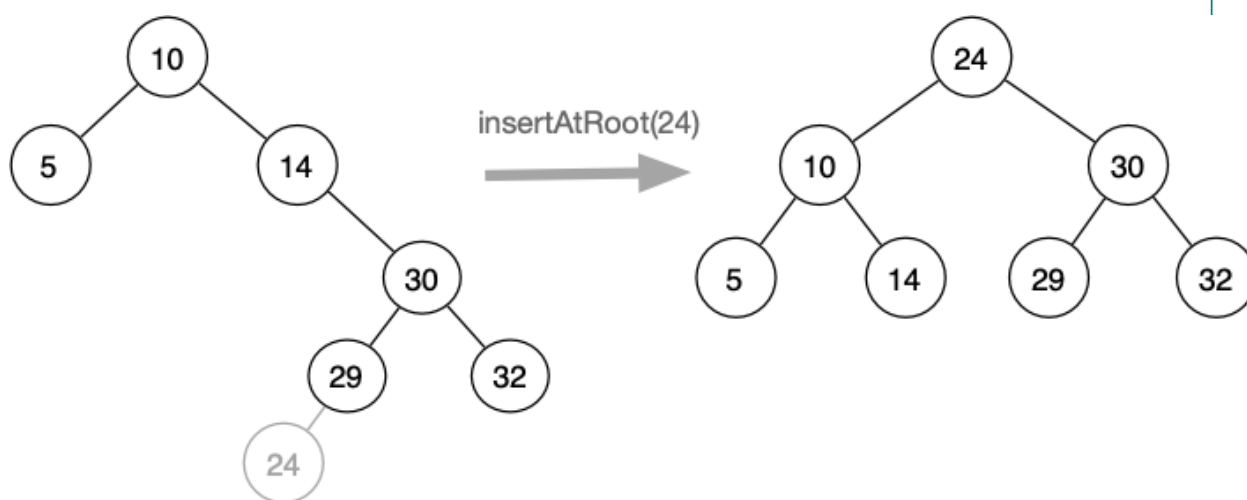
❖ ... Insertion at Root

Method for inserting at root:

- base case:
 - tree is empty; make new node and make it root
- recursive case:
 - insert new node as root of appropriate subtree
 - lift new node to root by rotation

❖ ... Insertion at Root

Example of inserting at root:



❖ ... Insertion at Root

Algorithm for inserting at root:

```
insertAtRoot(t, it):  
|   Input tree t, item it to be inserted  
|   Output modified tree with item at root  
|  
|   if t is empty tree then  
|       t = new node containing item  
|   else if item < root(t) then  
|       left(t) = insertAtRoot(left(t), it)  
|       t = rotateRight(t)  
|   else if it > root(t) then  
|       right(t) = insertAtRoot(right(t), it)  
|       t = rotateLeft(t)  
|   end if  
|   return t;
```

❖ ... Insertion at Root


Analysis of insertion-at-root:

- same complexity as for insertion-at-leaf: $O(\text{height})$
 - but cost is effectively doubled ... traverse down, rotate up
- tendency to be balanced, but no balance guarantee
- benefit comes in searching
 - for some applications, search favours recently-added items
 - insertion-at-root ensures these are close to root
- could even consider "move to root when found"
 - effectively provides "self-tuning" search tree

❖ Tree Partitioning

Tree partition operation **partition(tree, i)**

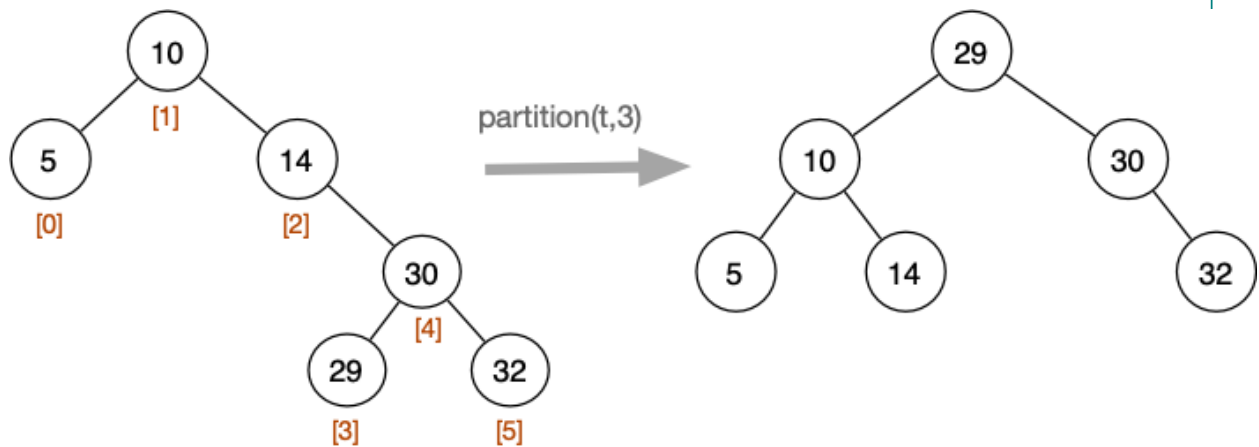
- re-arranges tree so that element with index i becomes root



For tree with N nodes, indices are $0..N-1$, in LNR order

❖ ... Tree Partitioning

Example of partition:



COMP2521 20T1 ♦ Balancing Search Trees ♦ [16/30]

❖ ... Tree Partitioning

Implementation of partition operation:

```
partition(tree,i):
|   Input   tree with n nodes, index i
|   Output tree with  $i^{\text{th}}$  item moved to the root
|
|   m=#nodes(left(tree))
|   if i < m then
|       left(tree)=partition(left(tree),i)
|       tree=rotateRight(tree)
|   else if i > m then
|       right(tree)=partition(right(tree),i-m-1)
|       tree=rotateLeft(tree)
|   end if
|   return tree
```

Note: $\text{size}(\text{tree}) = n$, $\text{size}(\text{left}(\text{tree})) = m$, $\text{size}(\text{right}(\text{tree})) = n - m - 1$

❖ ... Tree Partitioning

Analysis of tree partitioning

- no requirement for search (using element index instead)
- after each recursive partitioning step, one rotation
- overall cost similar to insert-at-root

Benefits

- tends to improve balance \Rightarrow improves search cost

❖ Periodic Rebalancing

An approach to maintaining balance:

- insert at leaves as before; periodically, rebalance the tree

```
| Input   tree, item  
| Output tree with item randomly inserted  
|  
| t=insertAtLeaf(tree,item)  
| if #nodes(t) mod k = 0 then  
|     t=rebalance(t)  
| end if  
| return t
```

When to rebalance? e.g. after every *k* insertions

❖ ... Periodic Rebalancing

A problem with this approach ...


- operation `#nodes()` has to traverse whole (sub)tree
- to improve efficiency, change node structure

```
typedef struct Node {  
    int  data;  
    int  nnodes;      // #nodes in my tree  
    Tree left, right; // subtrees  
} Node;
```

But maintaining `nnodes` requires extra work in other operations

❖ ... Periodic Rebalancing

How to rebalance a BST? Move median item to root.



COMP2521 20T1 ♦ Balancing Search Trees ♦ [21/30]

❖ ... Periodic Rebalancing

Implementation of rebalance:

```
rebalance(t):  
|   Input   tree t with n nodes  
|   Output  t rebalanced  
|  
|   if  $n \geq 3$  then  
|   |   // put node with median key at root  
|   |   t=partition(t,  $\lfloor n/2 \rfloor$ )  
|   |   // then rebalance each subtree  
|   |   left(t)=rebalance(left(t))  
|   |   right(t)=rebalance(right(t))  
|   end if  
|   return t
```

❖ ... Periodic Rebalancing

Analysis of rebalancing: visits every node $\Rightarrow O(N)$

Cost means not feasible to rebalance after each insertion.

When to rebalance? ... Some possibilities:

- after every k insertions
- whenever "imbalance" exceeds threshold

Either way, we tolerate worse search performance for periods of time.

Does it solve the problem? ... Not completely \Rightarrow Solution: real balanced trees (next week)

❖ Randomised BST Insertion

Reminder: order of insertion can dramatically affect shape of tree

Tree ADT has no control over order that keys are supplied.

We know that inserting in random order gives $O(\log_2 n)$ search

Can the algorithm itself introduce some randomness?

In the hope that this randomness helps to balance the tree ...

❖ ... Randomised BST Insertion

Approach: normally do leaf insert, randomly do root insert.

```
insertRandom(tree,item)
|   Input   tree, item
|   Output tree with item randomly inserted
|
|   if tree is empty then
|       return new node containing item
|   end if
|   // p/q chance of doing root insert
|   if random() mod q < p then
|       return insertAtRoot(tree,item)
|   else
|       return insertAtLeaf(tree,item)
|   end if
```

E.g. 30% chance \Rightarrow choose $p=3, q=10$

❖ ... Randomised BST Insertion

Cost analysis:

- similar to cost for inserting keys in random order:
 $O(\log_2 n)$
- does not rely on keys being supplied in random order

Approach can also be applied to deletion:

- standard method promotes inorder successor to root
- for the randomised method ...
 - promote inorder successor from right subtree, OR
 - promote inorder predecessor from left subtree

❖ An Application of BSTs: Sets

Trees provide efficient search.

Sets require efficient search

- to find where to insert/delete
- to test for set membership

Logical to implement a set ADT via binary search tree.

❖ ... An Application of BSTs: Sets

Assuming we have BST implementation with type **Tree**

- which precludes duplicate key values
- which implements insertion, search, deletion


then **Set** implementation is

- **SetInsert**(Set, Item) \equiv **TreeInsert**(Tree, Item)
- **SetDelete**(Set, Item) \equiv
 TreeDelete(Tree, Item.Key)
- **SetMember**(Set, Item) \equiv
 TreeSearch(Tree, Item.Key)


What about union? and intersection?

❖ ... An Application of BSTs: Sets


Sets implemented via Trees:



After SetInsert(s,4):



After SetDelete(s,2):



❖ ... An Application of BSTs: Sets

Concrete representation:

```
#include <Tree.h>

typedef struct SetRep {
    int    nelems;
    Tree   root;
} SetRep;

Set newSet() {
    Set S = malloc(sizeof(SetRep));
    assert(S != NULL);
    S->nelems = 0;
    S->root = newTree();
    return S;
}
```

Produced: 8 Jun 2020