

AVL Trees

- Better Balanced Binary Search Trees
- AVL Trees
- AVL Tree Examples
- AVL Insertion Algorithm
- Maintaining Balance/Height
- Searching AVL Trees
- Performance of AVL Trees

❖ Better Balanced Binary Search Trees

So far, we have seen ...

- randomised trees ... make poor performance unlikely
- occasional rebalance ... fix balance periodically
- splay trees ... reasonable amortized performance
- but all types still have $O(n)$ worst case

Ideally, we want both average/worst case to be $O(\log n)$

- AVL trees ... fix imbalances as soon as they occur
- 2-3-4 trees ... use varying-sized nodes to assist balance
- red-black trees ... isomorphic to 2-3-4, but binary nodes

❖ AVL Trees

Invented by Georgy Adelson-Velsky and Evgenii Landis (1962)

Goal:

- tree remains reasonably well-balanced
 $O(\log n)$
- cost of fixing imbalance is relatively cheap

Approach:

- insertion (at leaves) may cause imbalance
- repair balance as soon as we notice imbalance
- repairs done locally, not by overall tree restructure

❖ ... AVL Trees

A tree is unbalanced when $\text{abs}(\text{height}(\text{left}) - \text{height}(\text{right})) > 1$

This can be repaired by rotation:

- if left subtree too deep, rotate right
- if right subtree too deep, rotate left

Problem: determining height/depth of subtrees is expensive


- need to traverse whole subtree to find longest path

Solution: store balance data in each node (either height or balance)

- but extra effort needed to maintain this data on insertion

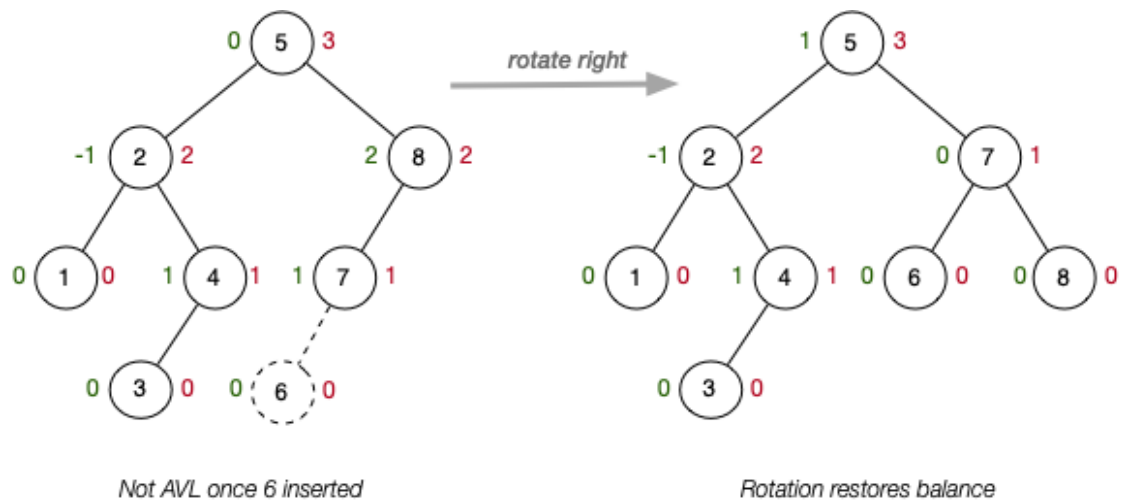
❖ AVL Tree Examples

Red numbers are height; green numbers are balance



❖ ... AVL Tree Examples

How an unbalanced tree can be rebalanced



COMP2521 20T2 ♦ AVL Trees [5/9]

❖ AVL Insertion Algorithm

Implementation of AVL insertion

```

insertAVL(tree,item):
  Input  tree, item
  Output tree with item AVL-inserted


  if tree is empty then
    return new node containing item
  else if item = data(tree) then
    return tree
  else
    if item < data(tree) then
      left(tree) = insertAVL(left(tree),item)
    else if item > data(tree) then
      right(tree) = insertAVL(right(tree),item)
    end if
    LHeight = height(left(tree))
    RHeight = height(right(tree))
    if (LHeight - RHeight) > 1 then
      if item > data(left(tree)) then
        left(tree) = rotateLeft(left(tree))
      end if
      tree=rotateRight(tree)
    else if (RHeight - LHeight) > 1 then
      if item < data(right(tree)) then
        right(tree) = rotateRight(right(tree))
      end if
      tree=rotateLeft(tree)
    end if
    return tree
  end if

```

❖ Maintaining Balance/Height


Store height in nodes; update on insertion;
compute balance

$$\text{balance} = \text{height}(\text{left}) - \text{height}(\text{right}) = 0 - 0 = 0$$



Leaves always have balance 0


$$\text{balance} = \text{height}(\text{left}) - \text{height}(\text{right}) = 1 - 0 = 1$$



If $\text{abs}(\text{balance}) > 1$ after updating, rebalance via rotation

❖ Searching AVL Trees

Exactly the same as for regular BSTs.




Height/balance measures are ignored

❖ Performance of AVL Trees

Analysis of AVL trees:

- trees are **height**-balanced; subtree depths differ by ± 1
- average/worst-case search performance of $O(\log n)$
- *require* extra data to be stored in each node (efficiency)
- *require* extra data to be maintained during insertion
- may not be **weight**-balanced; subtree sizes may differ



Produced: 15 Jun 2020